



Processing real time AIS messages and  
integrating with third party geo data using Sesam

Tom Bech

Sesam product developer  
[tom.bech@sesam.io](mailto:tom.bech@sesam.io)

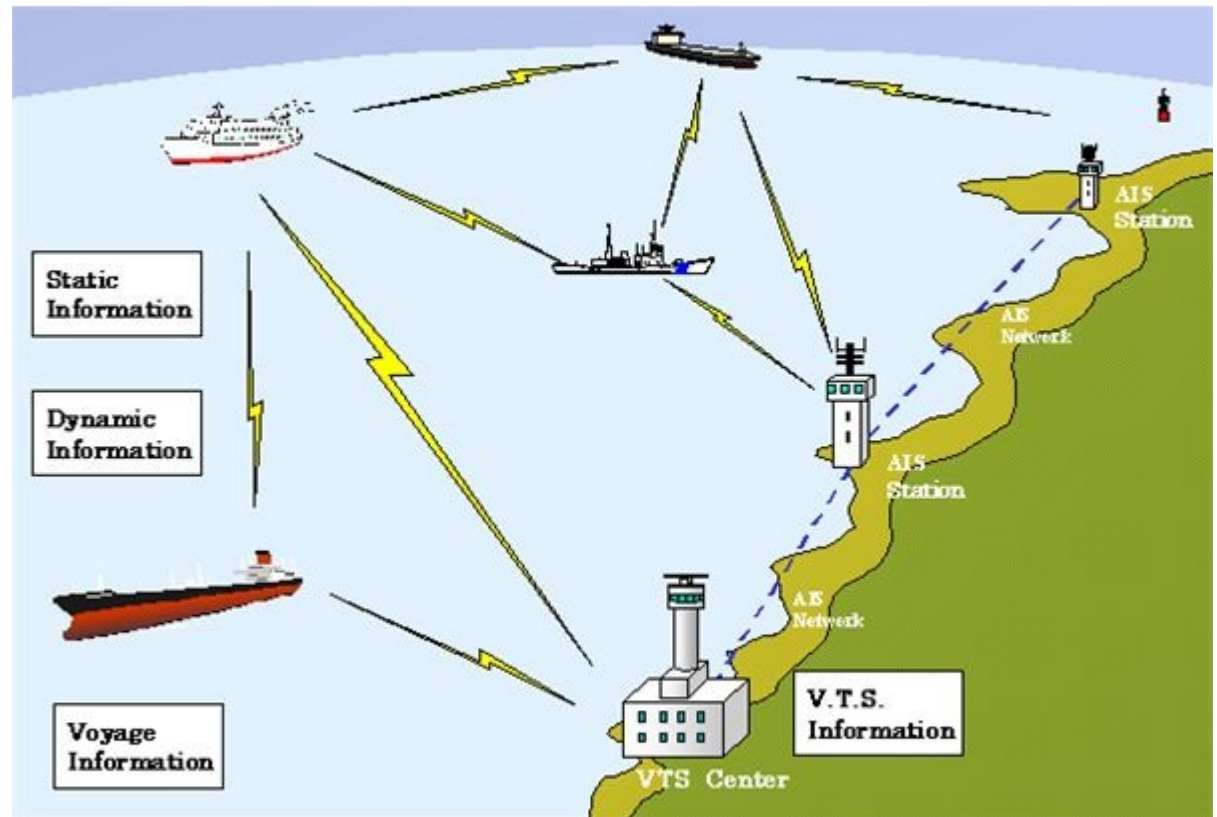
# Background & motivation

- Read interesting Norwegian Digi article about Gov't opening up “AIS” real time ship traffic data to the general public
- Disappointed to learn data not really open, article turn out to be about a web map *applicaton*
- Public *data* only available as socket with binary “raw” data (full access to “proper data” still subject to application & approval)

# AIS 101

- AIS = **A**utomatic *I*dentification **S**ystem
- UN standard (IEC 62320-1)
- Used world-wide for maritime traffic

- Real time messages from ships
  - ID, name and callsigns
  - Position updates, speed, bearing, status
  - Emergencies
  - Cargo information & destination
  - Metadata; type, size



# AIS con't

- Messages broadcast, rebroadcast
- Over-the-horizon media: VHF/UHF, satellites
  - Radar limited to visible range
- Picked up by transponders
  - Stations on land - i.e. basestations
  - Satelittes
- All larger vessels in world have AIS equipment
- Many leisure boats as well

# AIS in Norway

- “Kystverket” owns and operates
  - 50 base stations along the coast
  - Two satellites in polar orbit
- Infrastructure to process and analyze AIS data
  - Rebroadcasts AIS messages on the internet
  - ~30-100 messages/sec

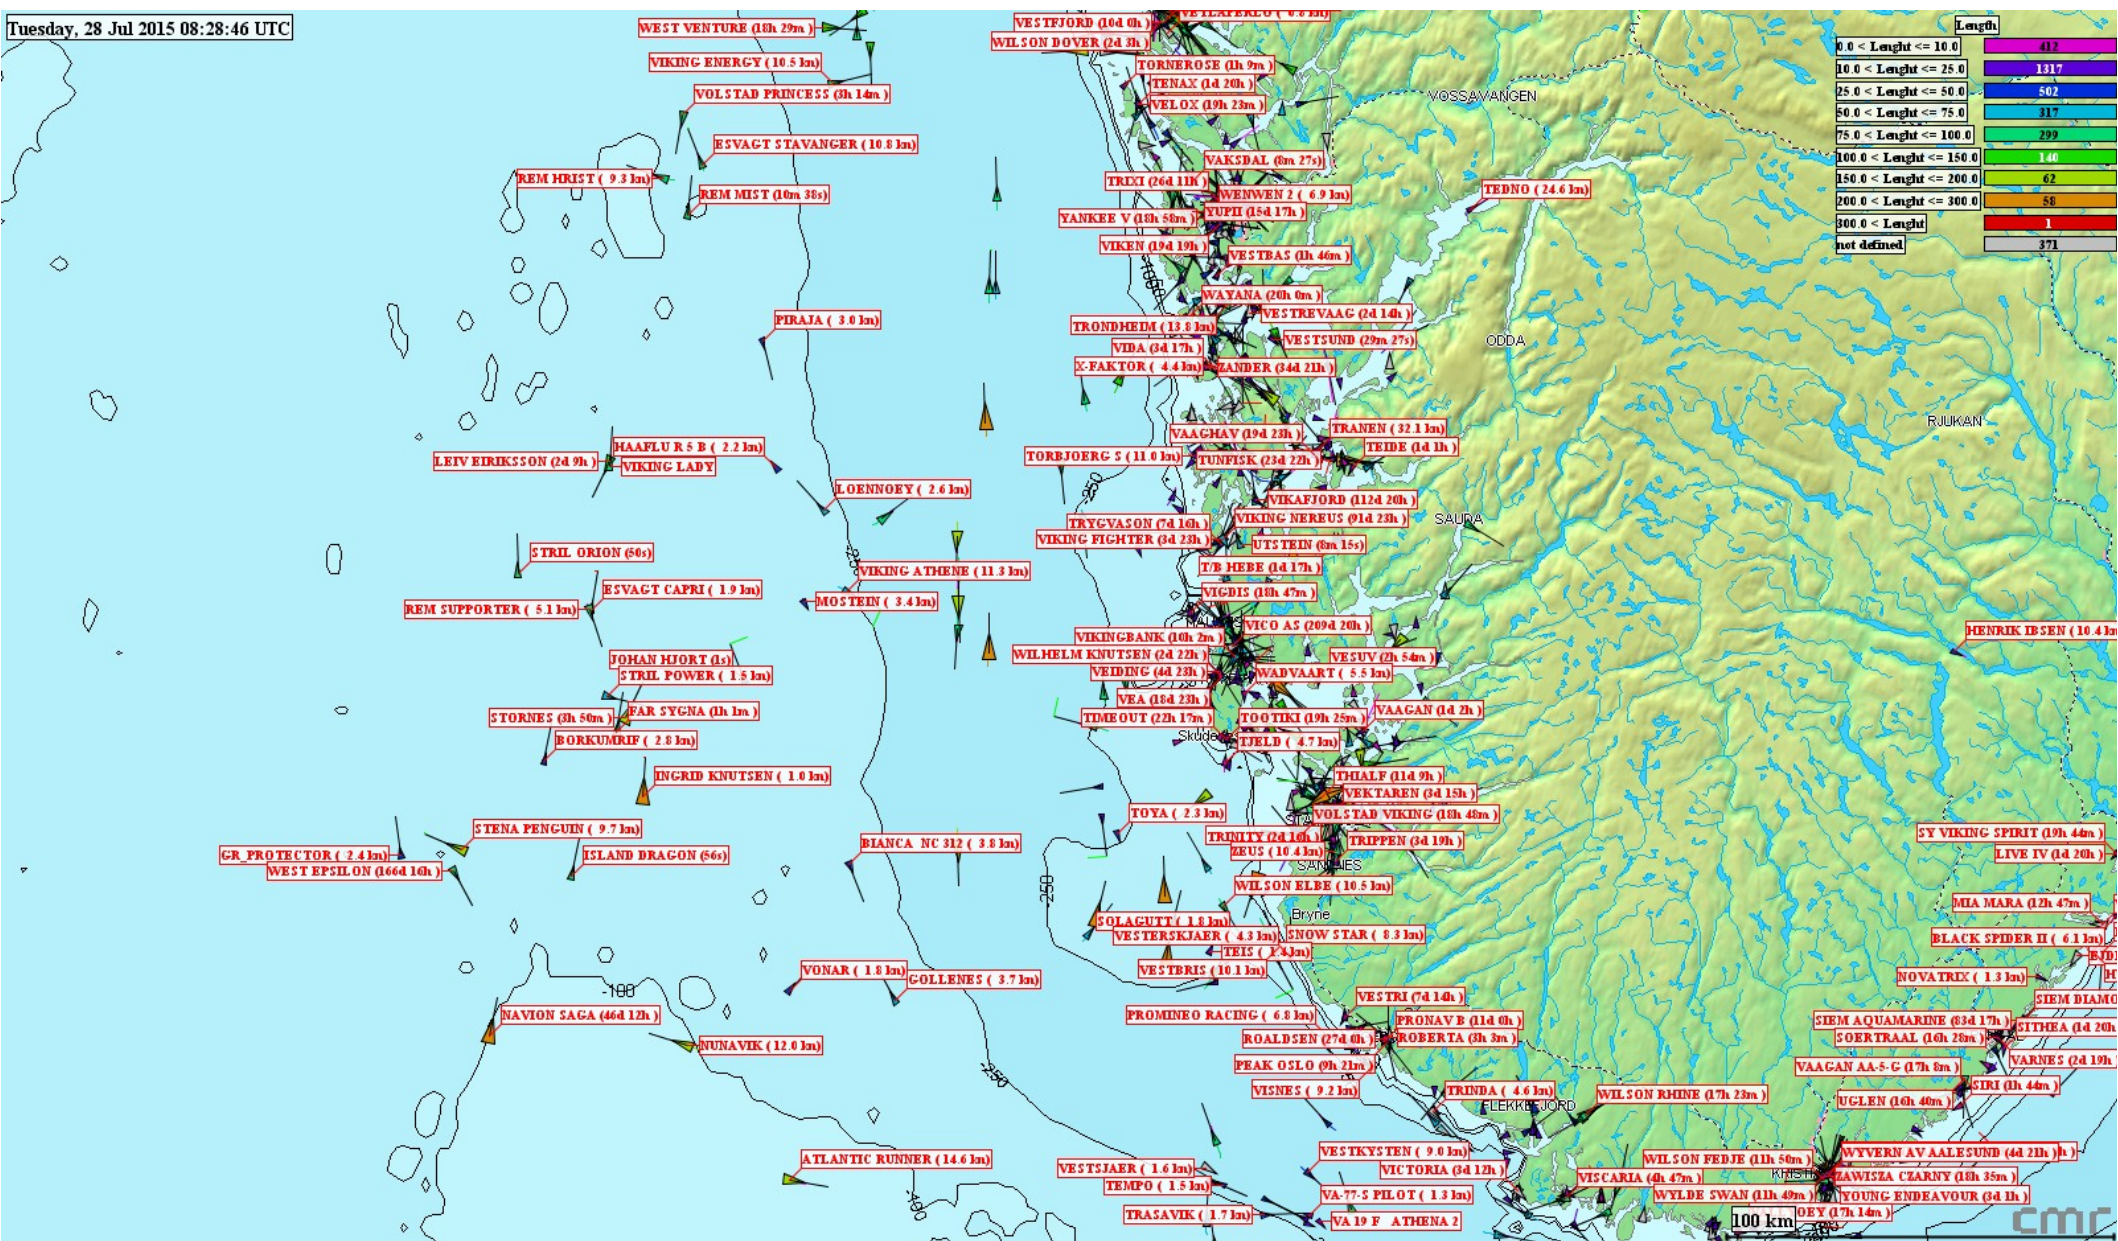


Friday, 08 May 2009 06:42:23 UTC





Tuesday, 28 Jul 2015 08:28:46 UTC





# The Challenge!

- Want to use AIS data + a search engine to do geo-queries like:
  - Which ships are near “Bergen”?
  - How many cargo ships are in Norwegian waters today?
  - Which ships are nearest vs another ship?
  - How many fishing boats are active near Lofoten at the moment?
- Additional challenge: from scratch - and in one day!

Programmer: Ooh. Very cool task, but pretty ambitious! Can I use Sesam as the integration platform?

Boss: Sure, use whatever tools you think is best for the job

# Challenge accepted!

# What is Sesam?

- Sesam is an integration platform
- Employs a Data Oriented Architecture (DOA) to make integration more manageable, scalable, powerful, expressive and agile
- Everything is data and everything is explicit
  - Including configuration, logging etc
- A datahub as a central component
  - Schema-less data store that consumes or is fed the data from the systems that are to be integrated
  - Solves the “cache-invalidation” problem in integration context
  - Push&pull source/transformation/sink + microservices = Pipe
  - Pipe+Pipe=Flow
  - Extensive REST API + command line client + GUI

# So..

- 1) Learn/research about AIS
- 2) Consume AIS data into Sesam somehow - write microservice?
- 3) Integrate AIS data with third party data (geolocated places) and compute nearest - another microservice?
- 4) Learn/research Elasticsearch
- 5) Insert result and changes into Elasticsearch
- 6) Figure how to do the queries
- 7) Profit!

..





# Starting point

- <http://www.kystverket.no/Nyheter/2016/september/apner-ais-for-publikum>
- A subpage mentions a AIS socket (IP+port)
- Tested it in Firefox. Got:

```
!BSVDM,1,1,,B,H3m<Od4N>F34u5K<=ojjn00H8220,0*73
!BSVDM,1,1,,A,35UeSP5000Puj;>V<B`;02mV0000,0*0F
!BSVDM,1,1,,A,13mM0u00001FIEdWnhUuL:OT0@NW,0*17
!BSVDM,2,1,1,A,53mN1J400000hluB2218E<=DF0T@5:1Di=@DTr0k0p?
154rdR2fLMevMeN88,0*73
!BSVDM,2,2,1,A,888888888880,2*3C
```

..



# The familiar first two hours

Google

 stackoverflow



 **GitHub**

  
WIKIPEDIA

# First round results

- An understanding of what AIS is in more detail
- Found docs for AIS spec in codemonkey-grokkable form:
  - <http://catb.org/gpsd/AIVDM.html>
- Found python lib for decoding AIS messages:
  - <https://github.com/schwehr/libais>
- Norwegian places with geo coords:
  - <http://www.erikbolstad.no/geo/noreg/postnummer>
- Nearest point search and geo calculus:
  - [https://en.wikipedia.org/wiki/K-d\\_tree](https://en.wikipedia.org/wiki/K-d_tree)
  - [https://en.wikipedia.org/wiki/Haversine\\_formula](https://en.wikipedia.org/wiki/Haversine_formula)
  - [https://en.wikipedia.org/wiki/Vincenty%27s\\_formulae](https://en.wikipedia.org/wiki/Vincenty%27s_formulae)

# Third and fourth hour



- Forked Sesam Python Microservice template on Github + wrote Python code to read and decode AIS stream using “libais”
- The usual loop; make stupid error. Fix it. Rinse & repeat.
- Configured JSON endpoint in Sesam & posted first AIS data to Sesam



# AIS MicroService Code

```
# Open socket to AIS service
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.connect((options.ais_server, int(options.ais_port)))

f = s.makefile()

with requests.session() as session:
    # This loop will not end as long as the socket is open
    for msg in ais.stream.decode(f):
        message = ais.compatibility.gpsd.mangle(msg)
        message["_id"] = "%s_%s" % (message["type"], message["mmsi"])

        json_data = json.dumps(message)
        resp = session.post(url, params={}, headers={"content-type": "application/json"},
                           data=json_data, verify=False, timeout=3600)

        resp.raise_for_status()
        resp.close()
```

# Sesam receiver conf

This pipe will set up a JSON receiver so we can use HTTP POST operations at

**`http://sesamservice:port/api/receivers/ais_data/entities`**

```
{  
  "_id": "ais_data",  
  "type": "pipe",  
  "source": {  
    "type": "http_endpoint"  
  }  
}
```

# AIS messages

```
{  
  "maneuver": 0,  
  "received_stations": 29,  
  "slot_timeout": 3,  
  "status": 0,  
  "second": 31,  
  "class": "AIS",  
  "scaled": True,  
  "course": 0,  
  "raim": True,  
  "type": 1,  
  "lat": 62.6781005859375,  
  "spare": 0,  
  "sync_state": 0,  
  "device": "stdin",  
  "repeat": 0,  
  "lon": 6.669294834136963,  
  "speed": 0,  
  "accuracy": True,  
  "status_text": "Under way using engine",  
  "turn": NaN,  
  "heading": 511,  
  "mmsi": 257389600  
}
```

Message type

```
{  
  "type": 24,  
  "device": "stdin",  
  "repeat": 0,  
  "shipname": "AAS KYSTSERVICE",  
  "mmsi": 257389600,  
  "class": "AIS",  
  "scaled": true,  
  "part_num": 0  
}
```

Unique ship ID

# Message deduplication

- AIS feed contain **LOTS** of duplicates (10:1)
  - Rebroadcasts + captured by multiple transponders
  - Static non-changing/repeating messages
  - Noise in data
- Sesam deduplicates automatically using hashing if given a unique entity ID
  - ID for message = MMSI + type
- Sesam keeps history of changes
  - Track ship positions over time!



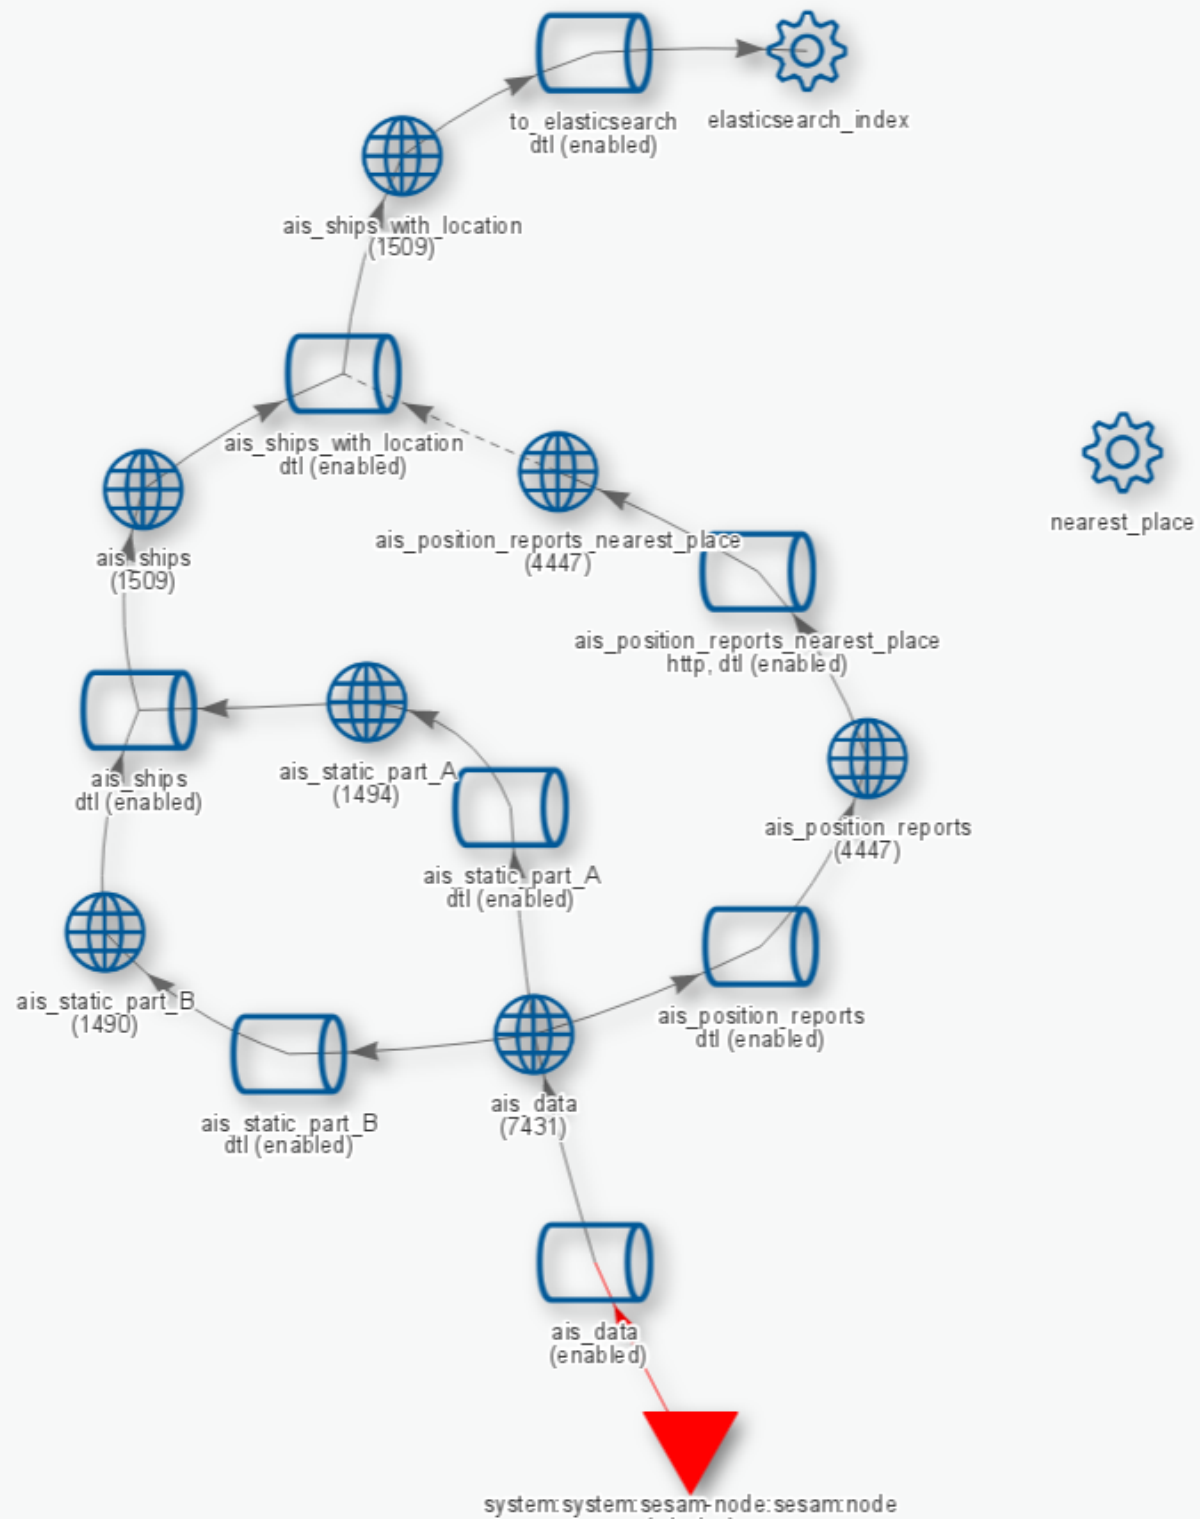
# Hours 4 and 5

## Splitting on message type

- Split static and dynamic messages on type
- Positional updates
  - AIS types 1-3 & 18-19
- Static metadata
  - Type 5 & 24
  - These may be multipart messages!

# Pipes and DTL

- Deduplicated AIS data from initial dataset piped to various new datasets based on type
  - Positional messages
  - Static data -> part A and B
  - Ship metadata from multipart messages (A U B)
  - Positional data enriched with nearest city info
  - Ship info with last known location
  - Final flow is outbound, i.e. Elasticsearch sink




# DTL transform for filtering on type

```
{
  "_id": "ais_position_reports",
  "type": "pipe",
  "source": {
    "type": "dataset",
    "dataset": "ais_data"
  },
  "transform": [
    {
      "type": "dtl",
      "name": "Filter out all but position reports (type 1-3 and 18-19)",
      "dataset": "ais_data",
      "transforms": {
        "default": [
          ["filter", ["or",
            ["eq", "_S.type", 1],
            ["eq", "_S.type", 2],
            ["eq", "_S.type", 3],
            ["eq", "_S.type", 18],
            ["eq", "_S.type", 19]]
          ],
          ["copy", "*"]
        ]
      }
    }
  ]
}
```


# Hours 6 and 7

## Nearest city microservice


- Sesam Python HTTP transform microservice template from github
- Inserts places into a KD-tree (BSP) on startup




**Iteration 1**



**Iteration 2**



**Iteration 3**



**Iteration 4**



# Nearest city microservice

- Given a set of entities, searches KD-tree for nearest city and inserts into entity
- Computes bearing and distance to city using Haversine- and Vincenty formula
- The enriched entity is returned to Sesam in response (JSON)

```
$ curl -s -XPOST 'http://localhost:5001/transform' -H "Content-type: application/json" -d '[{
  "_id": "jane", "name": "Jane Doe", "lat": 58.995903, "lon": 10.082722}]' | jq -S .
```

```
[
  {
    "_id": "jane",
    "lat": 58.995903,
    "lon": 10.082722,
    "name": "Jane Doe",
    "nearest_place": {
      "bearing": 170.2446453605429,
      "direction": "SSE",
      "distance": 3963.305464313277,
      "lat": 59.030966,
      "lon": 10.071019,
      "name": "Larvik",
      "postal_code": "3260"
    }
  }
]
```

Original entity

Added data

# 8 hour: Elasticsearch

- Googled how to run Elasticsearch:

```
docker pull elasticsearch
```

```
docker run --name elasticsearch -p 9200:9200 -p  
9300:9300 -d elasticsearch
```

- Googled how to configure a Elasticsearch index and wrote a schema for my entities

```
{
  "mappings": {
    "ship": {
      "properties": {
        "mmsi": {"type": "integer"},
        "callsign": {"type": "string"},
        "shipname": {"type": "string"},
        "length": {"type": "integer"},
        "width": {"type": "integer"},
        "position": {"type": "string"},
        "when": {"type": "date"},
        "vendor_id": {"type": "string"},
        "url": {"type": "string"},
        "location": {"type": "geo_point"}
      }
    }
  }
} => ships.json
```

**curl -XPUT http://<elasticsearch-ipaddress-here>:9200/ships @ships.json**

# Sesam Elasticsearch pipe

```
{
  "_id": "elasticsearch_index",
  "type": "system:elasticsearch",
  "hosts": ["172.17.0.2:9200"]
},
{
  "_id": "to_elasticsearch",
  "type": "pipe",
  "source": {
    "type": "dataset",
    "dataset": "ais_ships_with_location"
  },
  "sink": {
    "type": "elasticsearch",
    "system": "elasticsearch_index",
    "default_index": "ships",
    "default_type": "ship"
  }
}
```

# A sprinkle of DTL

- I also added a bit of DTL to massage the data *slightly* on the way out



```
"transform": [
{
  "type": "dtl",
  "name": "Transform to elasticsearch document",
  "dataset": "ais_ships_with_location",
  "transforms": {
    "default": [
      ["copy", "_id"],
      ["copy", "mmsi"],
      ["add", "length", ["+", "_S.to_stern", "_S.to_bow"]],
      ["add", "width", ["+", "_S.to_port", "_S.to_starboard"]],
      ["copy", "vendor_id"],
      ["copy", "callsign"],
      ["copy", "shipname"],
      ["copy", "url"],
      ["rename", "status_text", "status"],
      ["rename", "shiptype_text", "shiptype"],
      ["merge", ["apply", "apply-last-seen", "_S.last-seen-at"]]
    ],
    "apply-last-seen": [
      ["copy", "*"],
      ["add", "location", ["dict", ["list",
                                  ["list", "lat", "_S.lat"],
                                  ["list", "lon", "_S.lon"]
                                ]
                        ]],
      ["remove", "lat"],
      ["remove", "lon"]
    ]
  }
}]
```



# Moment of truth – queries!

- Ships near Bergen?

- `curl -XGET 'http://172.17.0.2:9200/ships/ship/_search?q=position:bergen*&pretty=true'`

- Fishingboats near Leknes (in Lofoten islands, northern Norway):

- `curl -XGET 'http://172.17.0.2:9200/ships/ship/_search?q=position:leknes*%20AND%20shiptype:fishing*&pretty=true'`

- All ships named something with "viking":

- `curl -XGET 'http://172.17.0.2:9200/ships/ship/_search?q=shipname:viking*&pretty=true'`

# Ships near a specific position

```
{
  "sort" : [
    {
      "_geo_distance" : {
        "location" : {
          "lat" : 59.902006,
          "lon" : 10.718077
        },
        "order" : "asc",
        "unit" : "km"
      }
    }
  ],
  "query": {
    "filtered" : {
      "query" : {
        "match_all" : {}
      },
      "filter" : {
        "geo_distance" : {
          "distance" : "5km",
          "location" : {
            "lat" : 59.902006,
            "lon" : 10.718077
          }
        }
      }
    }
  }
}
```

# Success!

Using Sesam, I was able to create a working (PoC) AIS solution  
- in a single day! :)



(Well, *technically* over one afternoon and subsequent morning – but in sum 8 hours. honest!)

# Show me the code!

- All code and full blog article

<https://github.com/sesam-io/ais-integration>

- Sign up for & download the Sesam appliance:

<https://beta.sesam.in/>